



**SYSTEMS
GOSSIP
MEETUP**

Learning large systems using peer-to-peer gossip

Policy Against Harassment at ACM Activities

<https://www.acm.org/about-acm/policy-against-harassment>

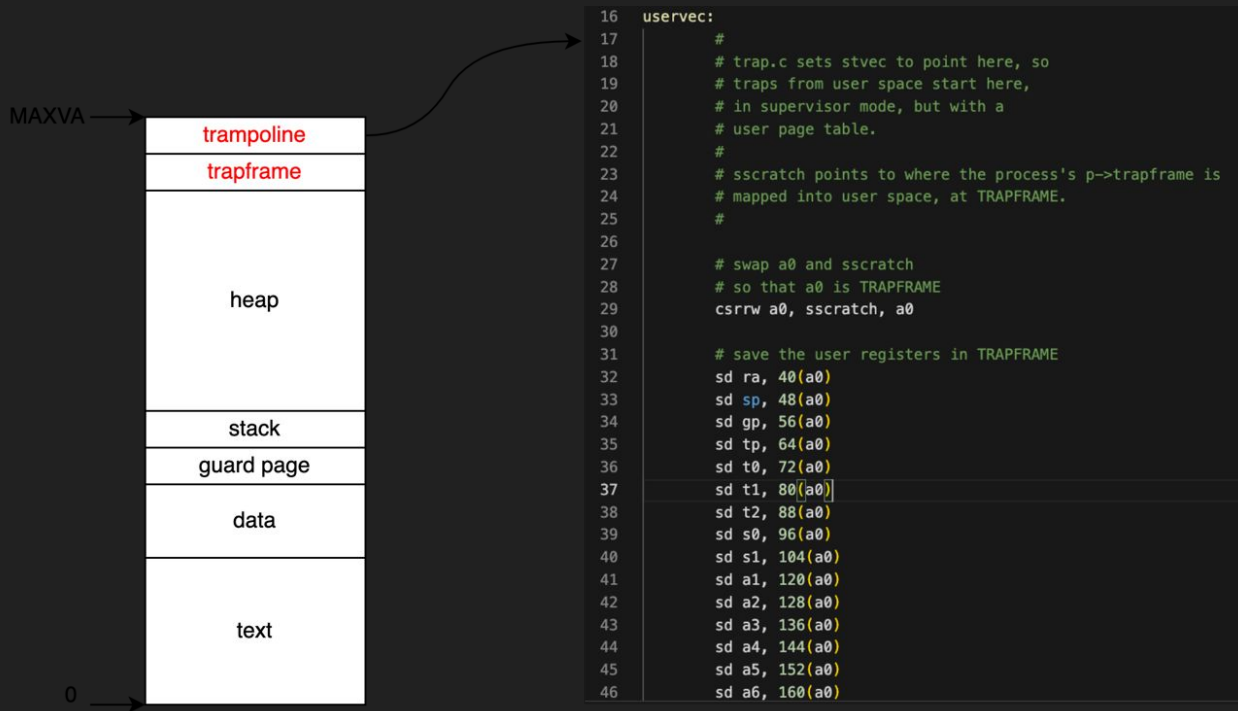
OS Meetup wants to encourage and preserve this open exchange of ideas, which requires an environment that enables all to participate without fear of personal harassment. We define harassment to include specific unacceptable factors and behaviors listed in the ACM's policy against harassment. Unacceptable behavior will not be tolerated.

Last Time

- spinlocks for multicore processing

Last Time

- spinlocks for multicore processing
- Syscall trap



Last Time

- spinlocks for multicore processing
- Syscall trap
- Timer interrupt trap

```
void
timerinit()
{
    // each CPU has a separate source of timer interrupts.
    int id = r_mhartid();

    // ask the CLINT for a timer interrupt.
    int interval = 1000000; // cycles; about 1/10th second in qemu.
    *(uint64*)CLINT_MTIMECMP(id) = *(uint64*)CLINT_MTIME + interval;

    // prepare information in scratch[] for timervec.
    // scratch[0..2] : space for timervec to save registers.
    // scratch[3] : address of CLINT MTIMECMP register.
    // scratch[4] : desired interval (in cycles) between timer interrupts.
    uint64 *scratch = &timer_scratch[id][0];
    scratch[3] = CLINT_MTIMECMP(id);
    scratch[4] = interval;
    w_mscratch((uint64)scratch);

    // set the machine-mode trap handler.
    w_mivec((uint64)timervec);

    // enable machine-mode interrupts.
    w_mstatus(r_mstatus() | MSTATUS_MIE);

    // enable machine-mode timer interrupts.
    w_mie(r_mie() | MIE_MTIE);
}
```

Today

- **Threads**
- **Context Switching in Xv6**
- **Scheduler**

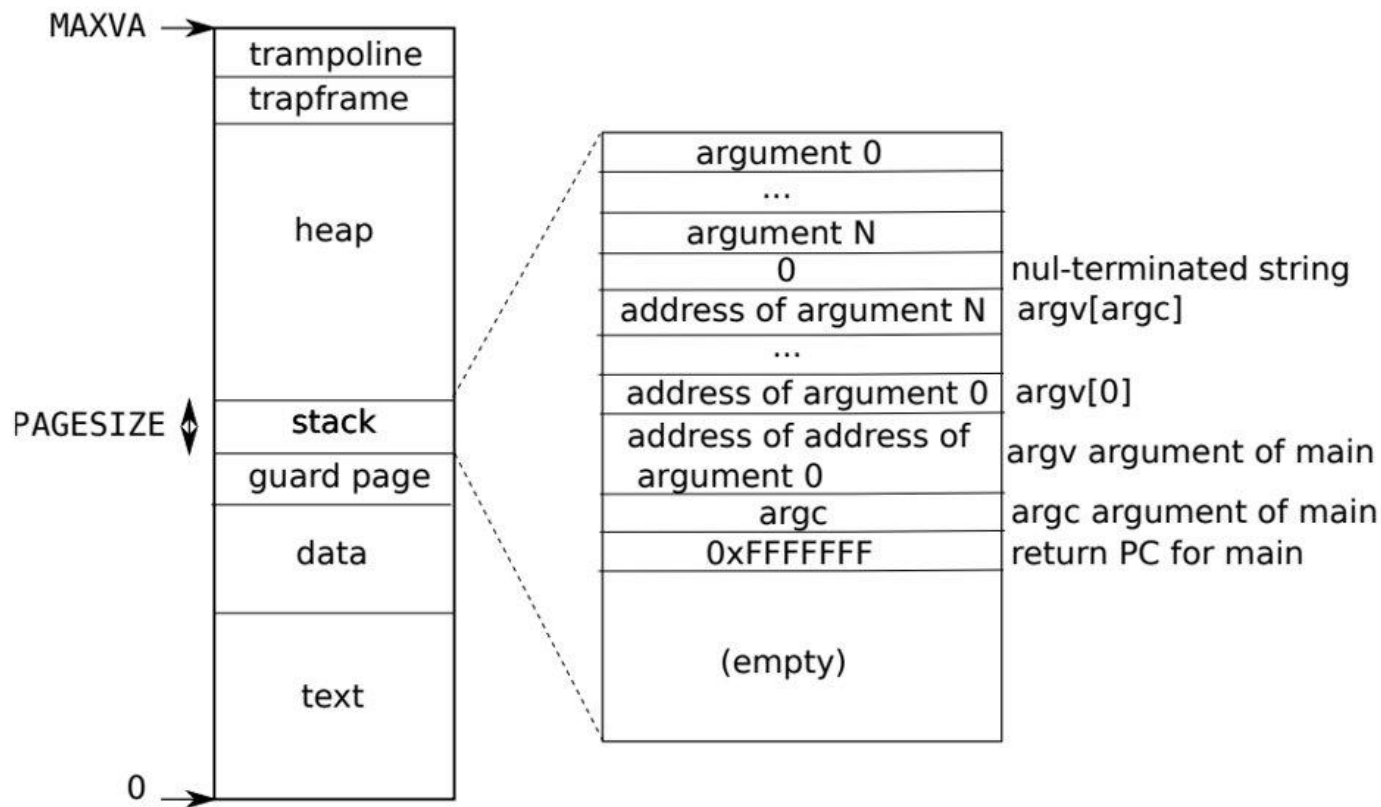


Figure 3.4: A process's user address space, with its initial stack.

Thread Switching

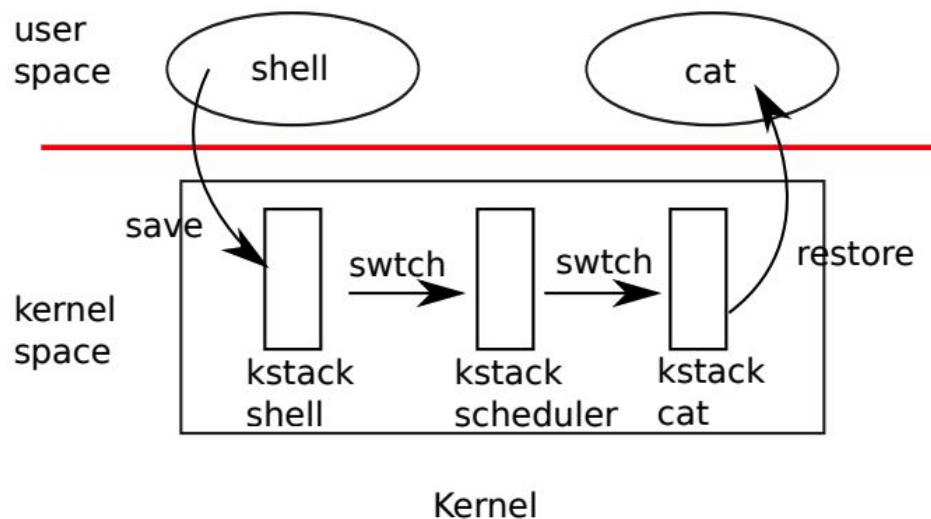


Figure 7.1: Switching from one user process to another. In this example, xv6 runs with one CPU (and thus one scheduler thread).

Threads

- In Xv6, each process has a thread of execution (we call it thread).
- If we want to support more processes than # of CPU cores, we want to **transparently** switch threads between processes.
- Each thread has its own thread stack, which saves local variables, parameters, return address.
- Each process has two types of stacks: user stack and kernel stack.
When the process runs in user space, it executes on user stack. When the process runs in kernel space, it executes on kernel stack.

Process Lifecycle Demo

```
// Per-process state
github-classroom[bot], 2 days ago | 1 author (github-classroom[bot])
struct proc {
    struct spinlock lock;

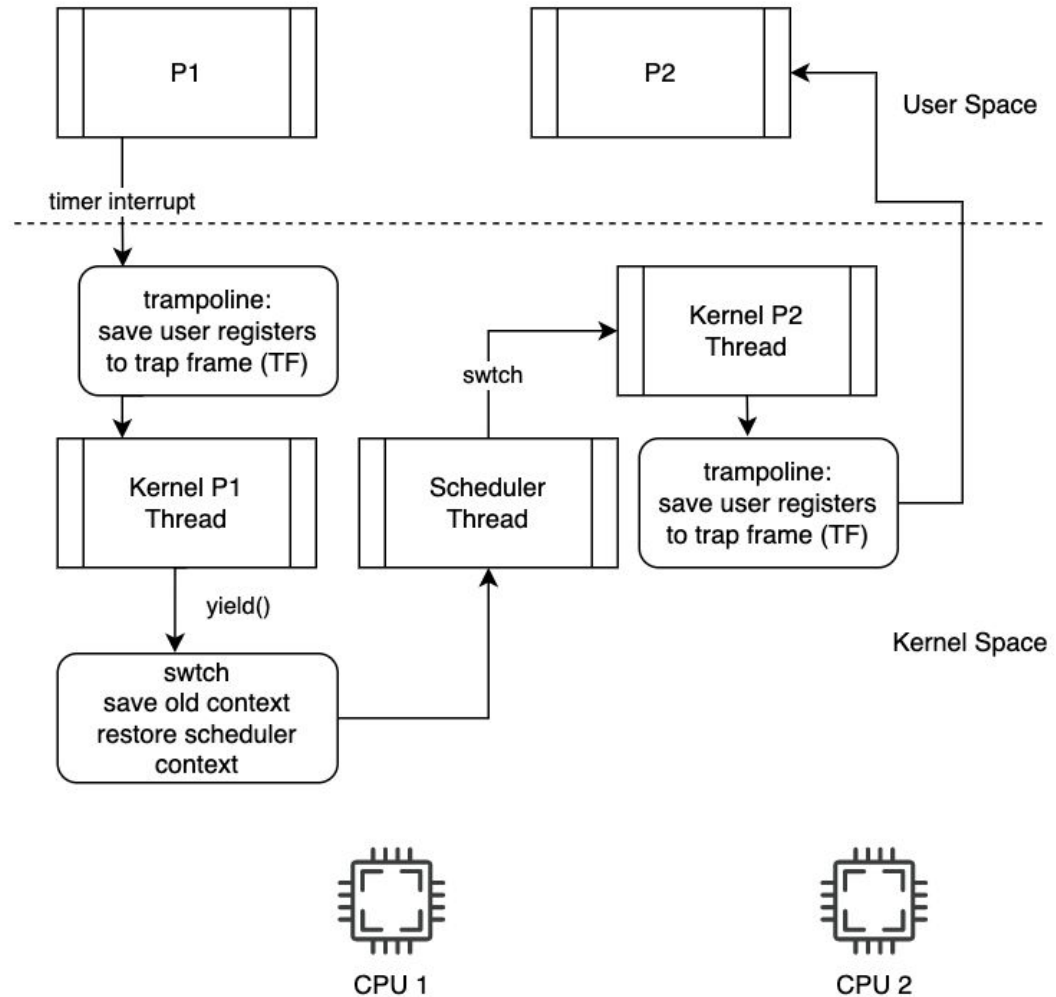
    // p->lock must be held when using these:
    enum procstate state;           // Process state
    void *chan;                     // If non-zero, sleeping on chan
    int killed;                     // If non-zero, have been killed
    int xstate;                     // Exit status to be returned to parent's wait
    int pid;                        // Process ID

    // wait_lock must be held when using this:
    struct proc *parent;            // Parent process

    // these are private to the process, so p->lock need not be held.
    uint64 kstack;                 // Virtual address of kernel stack
    uint64 sz;                      // Size of process memory (bytes)
    pagetable_t pagetable;        // User page table
    struct trapframe *trapframe;    // data page for trampoline.S
    struct context context;          // swtch() here to run process
    struct file *ofile[NOFILE];     // Open files
    struct inode *cwd;               // Current directory
    char name[16];                  // Process name (debugging)
};
```

```
enum procstate {
    UNUSED,
    USED,
    SLEEPING,
    RUNNABLE,
    RUNNING,
    ZOMBIE
};
```

Thread Switching



Next time

1. Lab 5 Cow is out! Due Oct 08th.
2. Lecture 12: Scheduling II