



**SYSTEMS  
GOSSIP  
MEETUP**

Learning large systems using peer-to-peer gossip

# Policy Against Harassment at ACM Activities

<https://www.acm.org/about-acm/policy-against-harassment>

OS Meetup wants to encourage and preserve this open exchange of ideas, which requires an environment that enables all to participate without fear of personal harassment. We define harassment to include specific unacceptable factors and behaviors listed in the ACM's policy against harassment. Unacceptable behavior will not be tolerated.

# Last Time

- Filesystem in Xv6
- Crash recovery in Xv6

# Last Time

- Filesystem in Xv6
- Crash recovery in Xv6

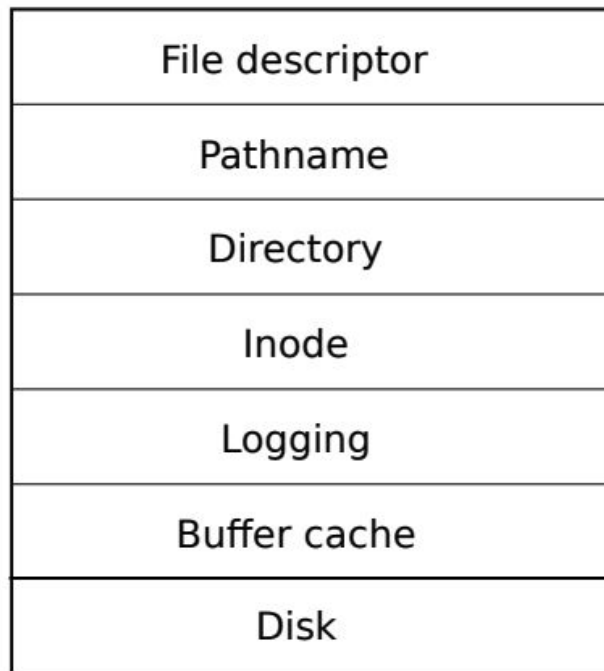


Figure 8.1: Layers of the xv6 file system.

# Last Time

- Filesystem in Xv6
- Crash recovery in Xv6




Figure 8.2: Structure of the xv6 file system.

# Today

- **Unix Ext3 Filesystem**

# Trade-off

- **Xv6 logging:**
  - Each system call is a transaction
  - At the end of the syscall
  - Write modified blocks to log on disk
  - Write blocks #s and done to log on disk
  - Write blocks to home location
  - Erase log
- **Immediately durable, but slow**
  - must commit after every syscall
  - must write to home location after commit

# File System Transactions

- Transaction is a series of operations that should be treated as one logical operation. Either all should be committed or none should be committed (All-or-Nothing)
- In Database, transactions are described using ACID
  - Atomicity
  - Consistency
  - Isolation
  - Durability



# Ext3

- Uses redo write-ahead logging for crash recovery
- In memory:
  - Block cache
  - Per-transaction bookkeeping
- On Disk:
  - Filesystem
  - Circular Log

# Ext3 Log

Super  
offset  
seq#

Tx4 Start  
LSN  
desc

Block  
with  
updated  
inode

Block  
with  
updated  
data  
bitmap

Data 1

Tx4 Cmt

Tx5 Start  
LSN  
desc

Data ...

# Ext3 Log

## Superblock

Offset  
Starting seq # of  
earliest valid  
transaction



## Tx N Start

Valid transaction

# Redo Write-ahead Logging

## Pre-crash:

- Write blocks that would be updated into the log
- Once all blocks are in the log, txn is committed. Ext3 issues an async writes to home blocks and metadata blocks

```
sys_unlink() {  
    h = start()  
    get(h, block #) ...  
    modify the block(s) in the cache  
    stop(h)  
}
```

# Redo Write-ahead Logging

## Post-crash:

- Scan the WAL , replay all the committed transactions and ignore uncommitted transactions.

```
sys_unlink() {  
    h = start()  
    get(h, block #) ...  
    modify the block(s) in the cache  
    stop(h)  
}
```

# Concurrent transactions

- Ext3 allows multiple transactions to concurrent execute
- Some fully committed transactions in the log
- Some doing log writes as part of the commit
- **One** open transaction thats accepting new syscalls.

# WAL Committing

- Block new syscalls
- Wait for in-progress syscalls to `stop()`
- Open a new transaction, unblock system calls
- Write descriptor to log on disk with a list of block numbers
- Write each block from cache to log on disk
- Wait
- Write the commit record
- Wait
- Async write cached blocks to go to their home locations

# WAL Committing

Why do we have to issue separately for Commit?

- The kernel will issue all of the log writes at once. The disk can reorder writes, which may cause issue if all records except the data record complete...
- During recovery phase, the filesystem will see that the above transaction has been committed, so a data block will permanently be lost.




# WAL Committing




# WAL Committing

Can be reused iff Tx1 committed and all blocks are  
written to home locations



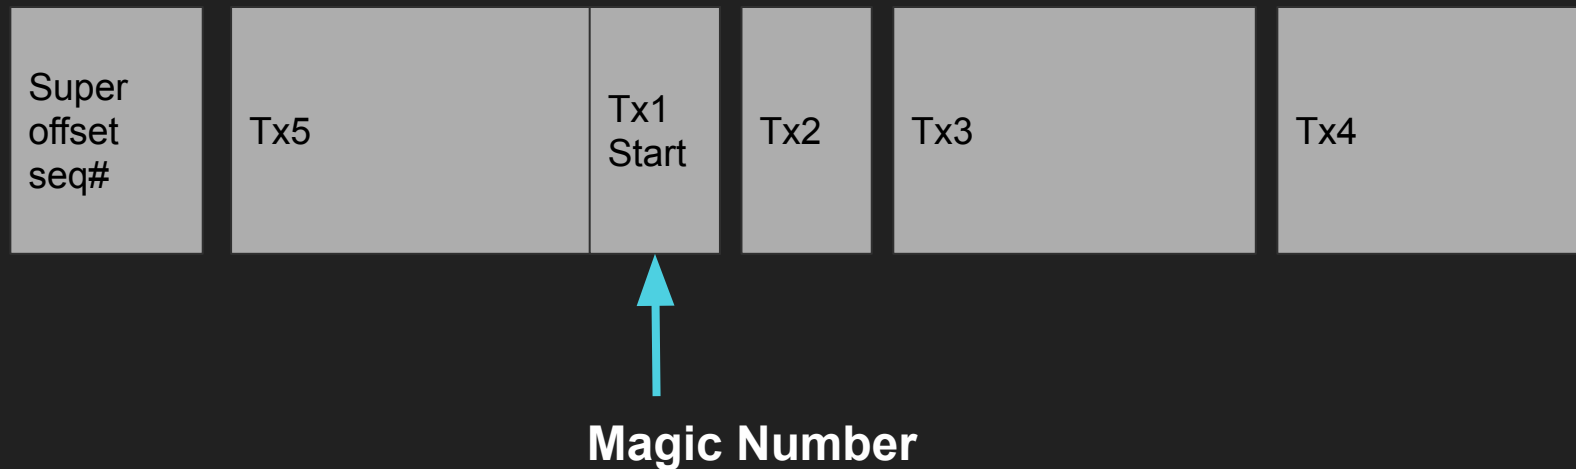
# Corner Case A

**What if block after last valid commit block looks like  
a log descriptor (Txn Start)?**




# Corner Case A

What if block after last valid commit block looks like  
a log descriptor (Txn Start)?



# Corner Case A

What if block after last valid commit block looks like  
a log descriptor (Txn Start)?



# Corner Case B

**What if block after last valid commit block looks like  
a Tx1 End?**



# Corner Case B

What if block after last valid commit block looks like  
a Tx1 End?



Log Sequence Number

# Summary of ext3 rules (lecture)

- Don't write metadata block to on-disk FS until committed in on-disk log
- Don't overwrite a block in buffer cache before it's in the log
- Don't free on-disk log transaction until all blocks have been written to FS



# Next time

1. Lab 6 Threads is out! Due Nov 12th.
2. Guest Talk!