



**SYSTEMS  
GOSSIP  
MEETUP**

**Learning large systems using peer-to-peer gossip**

# Policy Against Harassment at ACM Activities

<https://www.acm.org/about-acm/policy-against-harassment>

OS Meetup wants to encourage and preserve this open exchange of ideas, which requires an environment that enables all to participate without fear of personal harassment. We define harassment to include specific unacceptable factors and behaviors listed in the ACM's policy against harassment. Unacceptable behavior will not be tolerated.

# L4 Microkernels

# Today

- What is microkernel?
- Why is it important?
- Microkernel Design
- Microkernel Implementation: L4
- Conclusion

# What is Microkernel

- Move most operating system kernel functionality to user-space processes
- Minimal kernel
  - Address space
  - Threads
  - IPC (inter-process communication)
- 1980s hot research topic
- Examples: CMU's Mach, L4

## TOWARD REAL MICROKERNELS

*The inefficient, inflexible first generation inspired development of the vastly improved second generation, which may yet support a variety of operating systems.*



THE microkernel story is full of good ideas and blind alleys. The story began with enthusiasm about the promised dramatic increase in flexibility, safety, and modularity. But over the years, enthusiasm changed to disappointment, because the first-generation microkernels were inefficient and inflexible. • Today, we observe radically new approaches to the microkernel idea that seek to avoid the old mistakes while overcoming the old constraints on flexibility and performance. The second-generation microkernels may be a basis for all types of operating systems, including timesharing, multimedia, and soft and hard real time.

### The Kernel Vision

Traditionally, the word kernel denotes the mandatory part of the operating system common to all other software. The kernel can use all features of a processor (e.g., programming the memory management unit); software running in user mode cannot execute such safety-critical operations.

Most early operating systems were implemented by

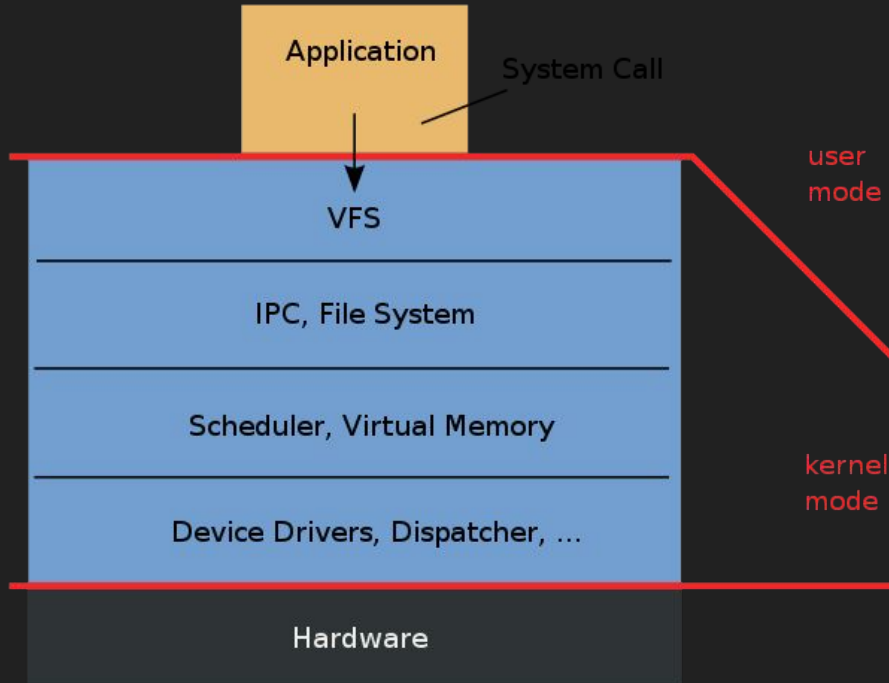
paging, and more—was packed into a single kernel.

In contrast, the microkernel approach involves minimizing the kernel and implementing servers outside the kernel. Ideally, the kernel implements only address spaces, interprocess communication (IPC), and basic scheduling. All servers—even device drivers—run in user mode and are treated exactly like any other application by the kernel. Since each serv-

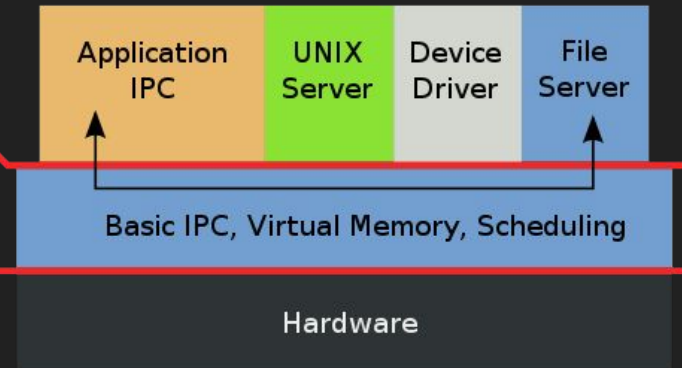
# Why is it important?

- **Elegant and clean**
  - **Minimal trusted computing base**
  - **Smaller attack surface**
- **Easy to prove**
- **More secure**
- **Easy to optimize**
- **More robust and extensible**
- **Inspires VM designs**

## Monolithic Kernel based Operating System



## Microkernel based Operating System



# Microkernel Evolution

## First gen (Mach '87)

- Bad performance
- 180 syscalls
- 100 kLOC
- 100 us IPC

### Improving IPC by Kernel Design

Jochen Liedtke

Proceeding of the 14<sup>th</sup> ACM Symposium on Operating  
Systems Principles

Asheville, North Carolina

1993





# Microkernel Evolution

## First gen (Mach '87)

- Bad performance
- 180 syscalls
- 100 kLOC
- 100 us IPC

## Sec gen (L4 '95)

- Great perf
- 7 syscalls
- 10 kLOC
- 1 us IPC
- Used in iPhone,  
iPad, Mac etc.

## Third gen (seL4 '09)

- Verifiable
- 3 syscalls
- 9 kLOC
- 0.1 us IPC

<https://support.apple.com/guide/security/secure-enclave-sec59b0b31ff/web>  
[https://trustworthy.systems/publications/nicta\\_full\\_text/8988.pdf](https://trustworthy.systems/publications/nicta_full_text/8988.pdf)

# L4 Basic Abstraction

## Task

Threads: each has a  
global ID

Own address space

## Message

From thread ID  
To thread ID  
Data

## Microkernel L4

Manage tasks, context switch, IPC (sending messages  
between tasks)

# L4 System Calls

## 1. ipc()

- a. Message passing, combining sending and receiving operations

## 2. fpage\_unmap()

- a. Revoke mappings

## 3. task\_new()

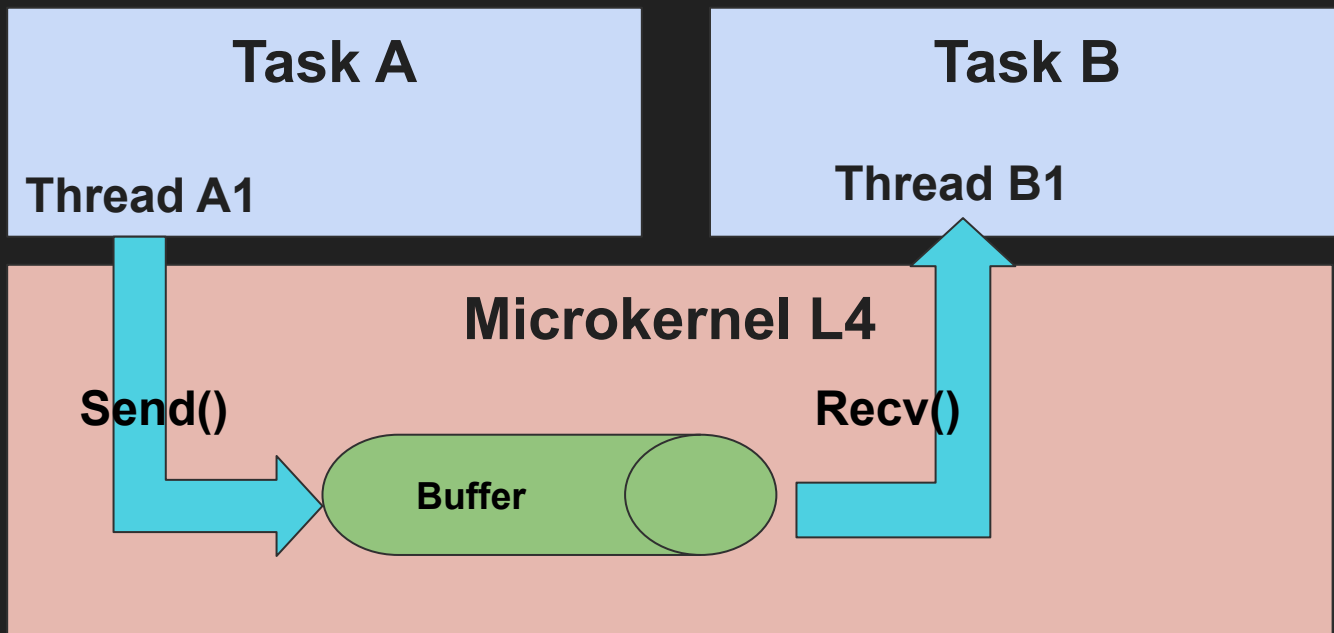
- a. create/delete task/address space

## 4. lthread\_ex\_regs()

- a. create/manipulate thread

# Problem: IPC Performance

Observation: microkernel programs do lots of IPC!

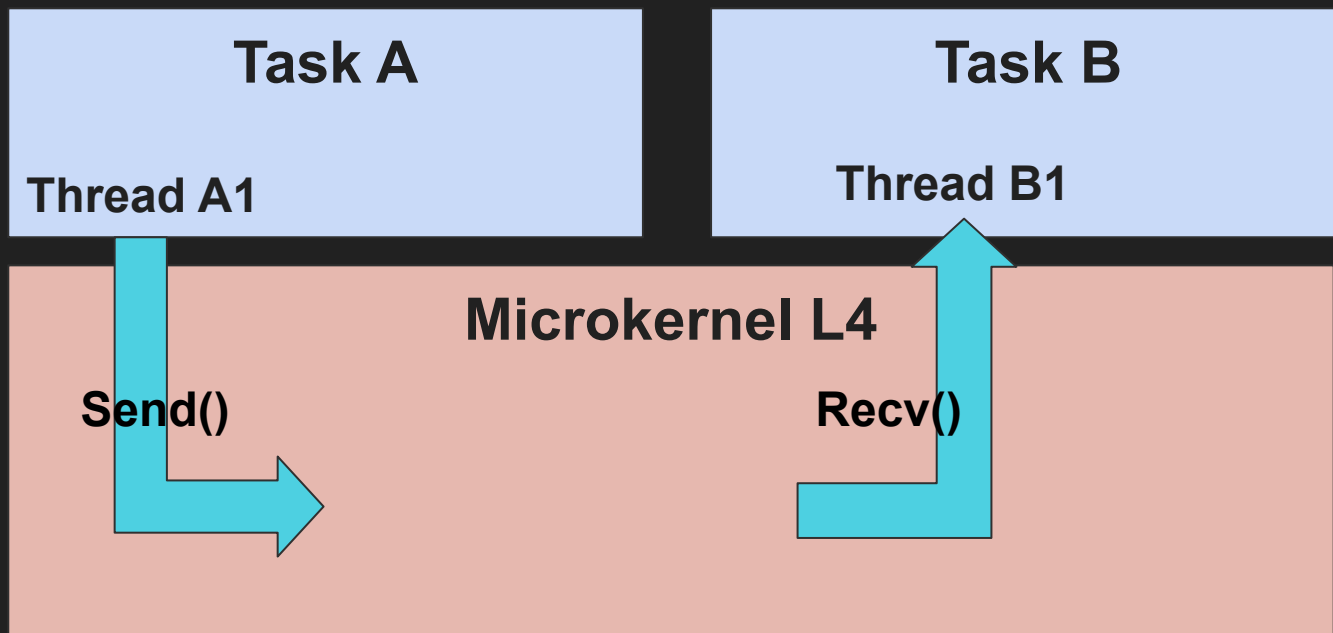


# Problem: IPC Performance

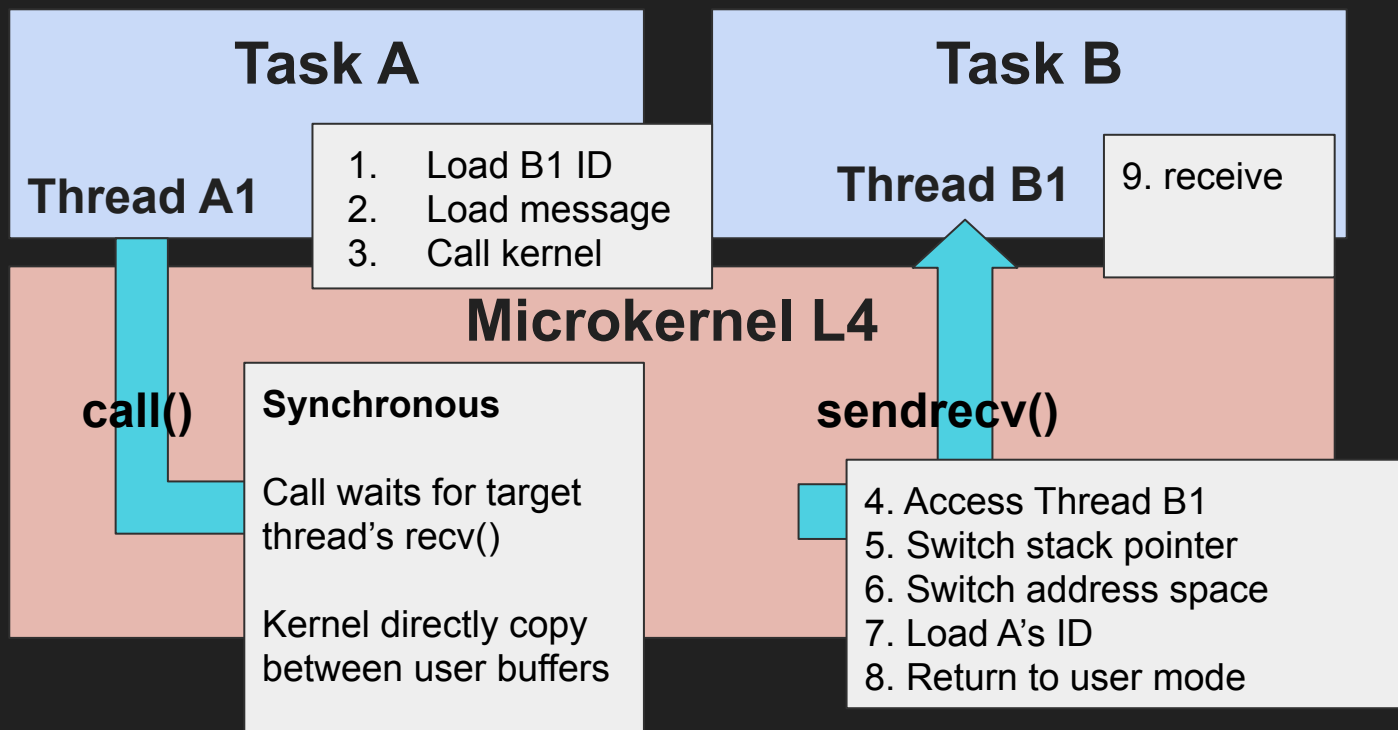
With the usual request-response pattern (RPC)

1. 4 system calls
  - a. `send()` -> `recv()`
  - b. `recv()` -> `send()`
2. Each system call may disturb CPU's caches
3. Four message copies (two for request, two for reply)
4. Two context switches, two schedulings

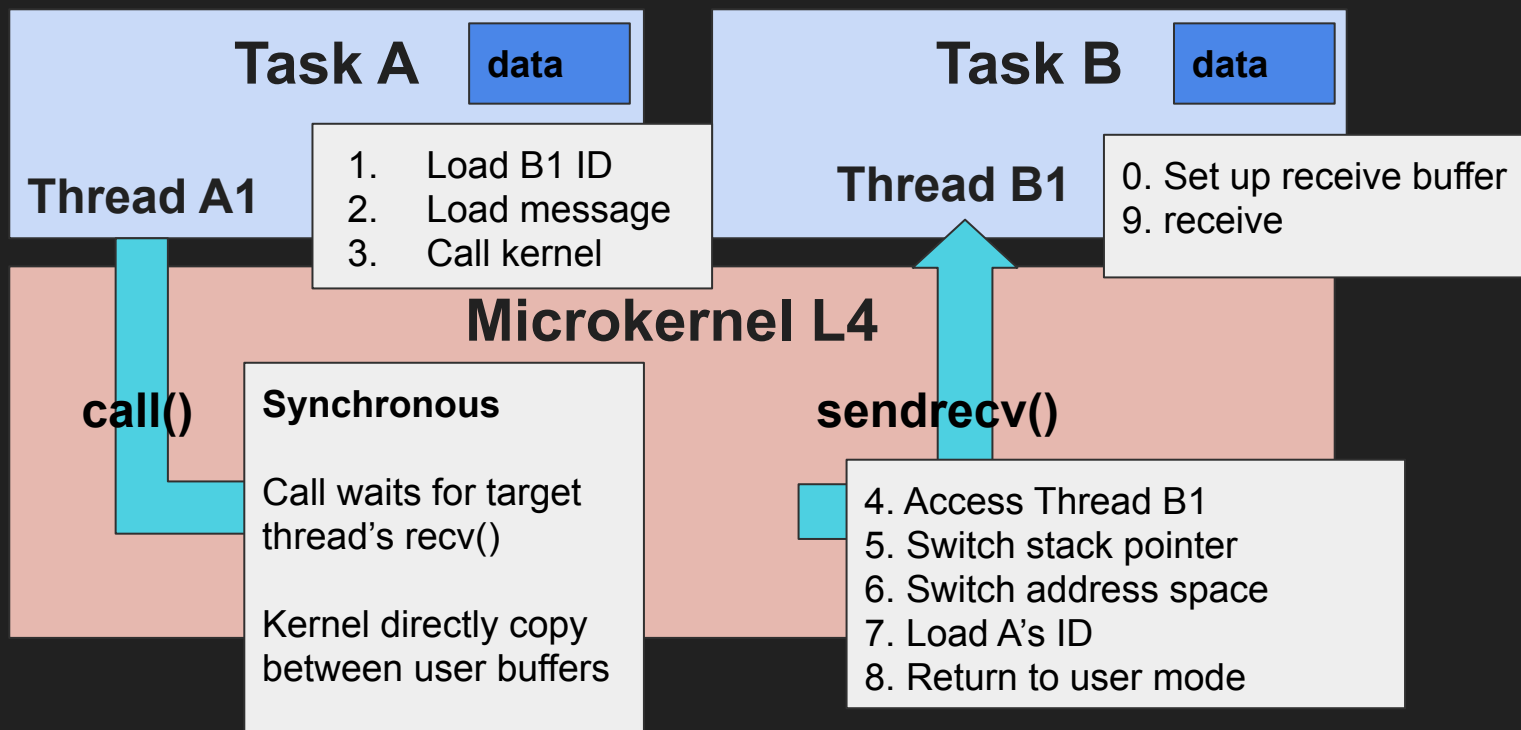
# Improving IPC by Kernel Design



# Improving IPC by Kernel Design



# Improving IPC by Kernel Design





# Improving IPC by Kernel Design

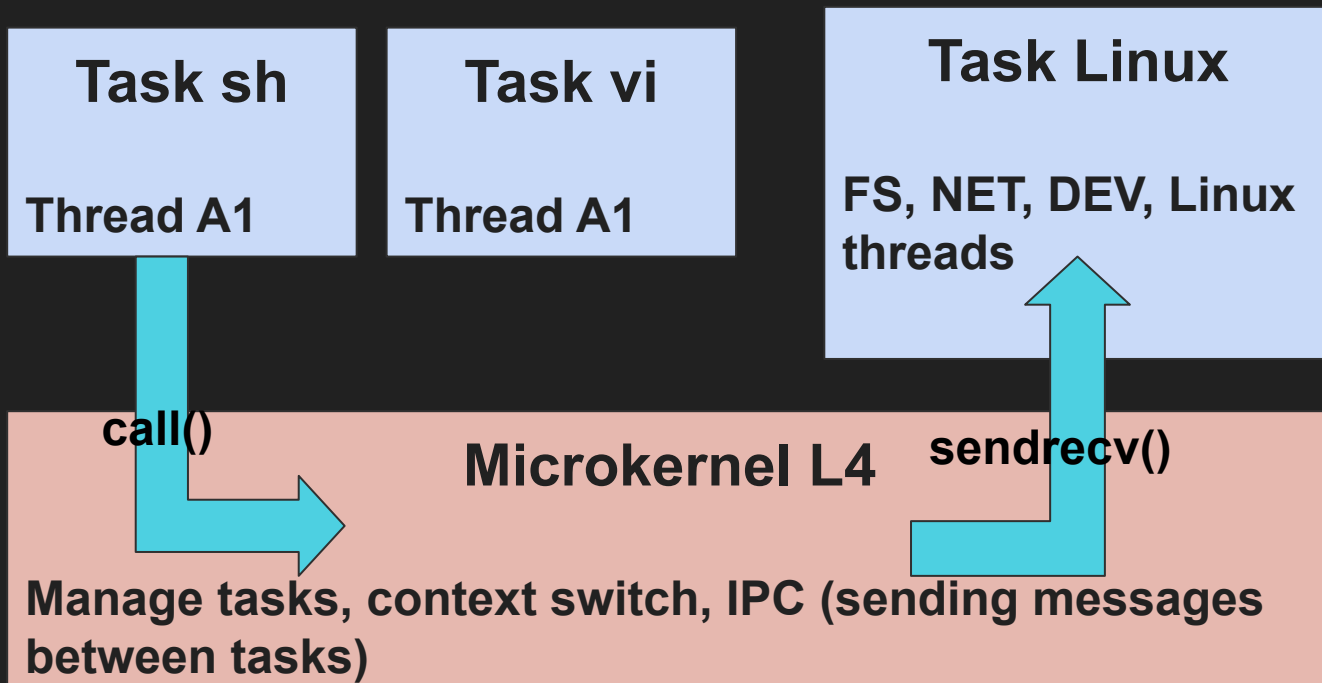
**Result: 20x reduction in IPC cost**

**2x reduction in user/kernel crossings**

**Makes microkernel performance comparable to monolithic kernel**

# Today's paper

## The performance of microkernel-based systems



# Today's paper

System	Time	Cycles
Linux	1.68 $\mu$ s	223
L <sup>4</sup> Linux	3.95 $\mu$ s	526
L <sup>4</sup> Linux (trampoline)	5.66 $\mu$ s	753
MkLinux in-kernel	15.41 $\mu$ s	2050
MkLinux user	110.60 $\mu$ s	14710

Table 2: *getpid* system-call costs on the different implementations. (133 MHz Pentium)




Figure 8: *AIM Multiuser Benchmark Suite VII*. Real time per benchmark run depending on AIM load units. (133 MHz Pentium)

# Conclusion

**What's current situation for microkernel?**

- 1. Used often in embedded computing. Apple “enclave” processor.**
- 2. Never caught on for general computing**
- 3. The OS server idea inspires Virtual Machines (VM).**

# Reference

1. <https://www.youtube.com/watch?v=dM9PLdaTpnA>
2. <https://www.cs.ubc.ca/~norm/508/2007W1/summaries/paper15/index.html>
3. [https://www.youtube.com/watch?v=mRr1lCJse\\_I](https://www.youtube.com/watch?v=mRr1lCJse_I)
4. <https://www.youtube.com/watch?v=oPjwyykO1bl>
5. <https://pdos.csail.mit.edu/6.S081/2021/lec/l-organization.txt>
6. [https://trustworthy.systems/publications/nicta\\_full\\_text/8988.pdf](https://trustworthy.systems/publications/nicta_full_text/8988.pdf)
7. <https://dl.acm.org/doi/10.1145/168619.168633>