

Read Copy Update

MIT 6.S081 Lec 23
Michael

Agenda

- Problem Motivation
- Read Write Lock
- Read Copy Update

Workload in Question

When:

- **write does happen but infrequently**
 - mostly read, rarely updated linked list accessible by kernel code
 - read the data, do a lot of computation and finally publish the result (scientific computing)
 - dynamically growing container: lots of read, but infrequently resizing

Since:

- we have many cores and we want to keep these cores busy

We need:

- concurrency control

Approaches



Performance Comparison

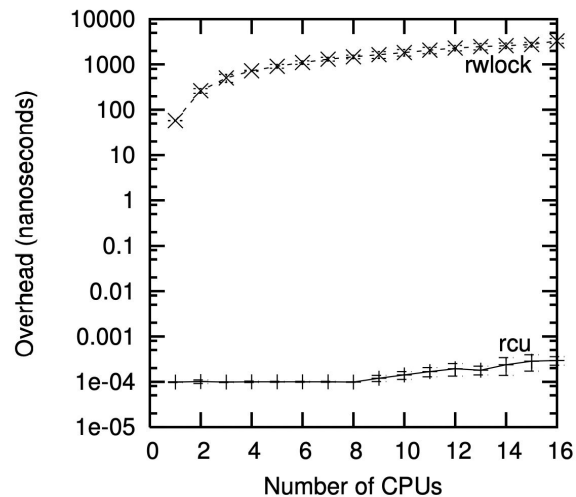


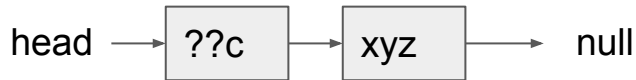
Figure 8: The overhead of entering and completing an RCU critical section, and acquiring and releasing a read-write lock.

Linked List: What can go wrong without synchronization

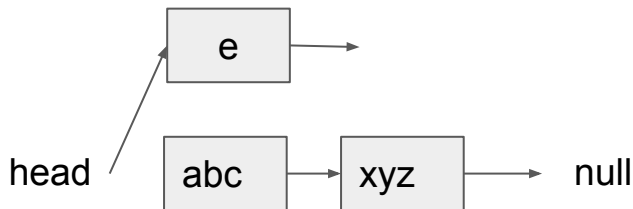


in kernel, many readers one writer, without sync:

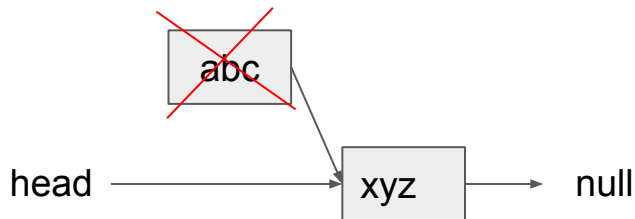
1. modify:



2. add node:



3. delete node:



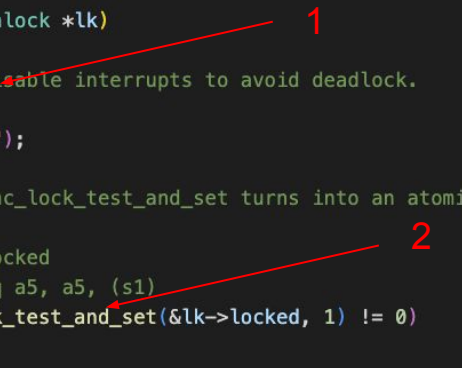
Spinlock

```
// Acquire the lock.
// Loops (spins) until the lock is acquired.
void
acquire(struct spinlock *lk)
{
    push_off(); // disable interrupts to avoid deadlock.
    if(holding(lk))
        panic("acquire");

    // On RISC-V, sync_lock_test_and_set turns into an atomic swap:
    //  a5 = 1
    //  s1 = &lk->locked
    //  amoswap.w.aq a5, a5, (s1)
    while(__sync_lock_test_and_set(&lk->locked, 1) != 0)
        ;

    // Tell the C compiler and the processor to not move loads or stores
    // past this point, to ensure that the critical section's memory
    // references happen strictly after the lock is acquired.
    // On RISC-V, this emits a fence instruction.
    __sync_synchronize();

    // Record info about lock acquisition for holding() and debugging.
    lk->cpu = mycpu();
}
```



1

2

```
// Release the lock.
void
release(struct spinlock *lk)
{
    if(!holding(lk))
        panic("release");

    lk->cpu = 0;

    // Tell the C compiler and the CPU to not move loads or stores
    // past this point, to ensure that all the stores in the critical
    // section are visible to other CPUs before the lock is released,
    // and that loads in the critical section occur strictly before
    // the lock is released.
    // On RISC-V, this emits a fence instruction.
    __sync_synchronize();

    // Release the lock, equivalent to lk->locked = 0.
    // This code doesn't use a C assignment, since the C standard
    // implies that an assignment might be implemented with
    // multiple store instructions.
    // On RISC-V, sync_lock_release turns into an atomic swap:
    //  s1 = &lk->locked
    //  amoswap.w zero, zero, (s1)
    __sync_lock_release(&lk->locked);

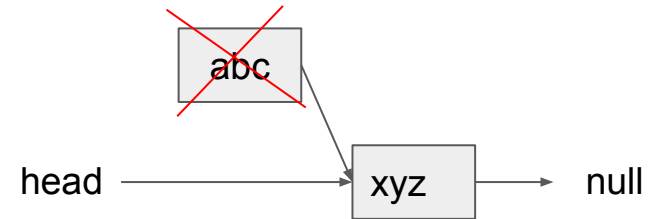
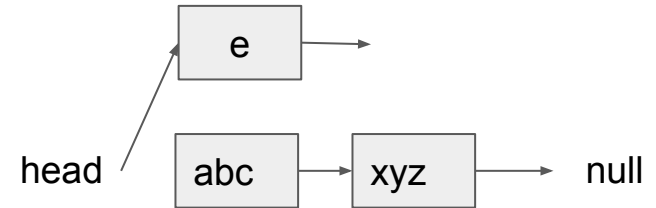
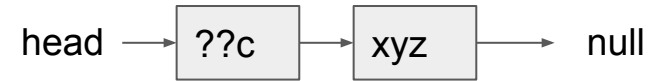
    pop_off();
}
```

RWLock: API

- Reader:
 - `reader_lock(l)`
 - `reader_unlock(l)`
- Writer:
 - `writer_lock(l)`
 - `writer_unlock(l)`

RWLock: Implementation

```
//  
// A simplified version of Linux's read/write lock.  
//  
// n=0  -> not locked  
// n=-1 -> locked by one writer  
// n>0  -> locked by n readers  
struct rwlock {  
    int n;  
};  
  
r_lock(l):  
    while 1:  
        x = l->n  
        if x < 0  
            continue  
        if CAS(&l->n, x, x + 1)  
            return  
  
// CAS(p, a, b) is atomic compare-and-swap instruction  
//   if *p == a, set *p = b, return true  
//   else return false  
  
w_lock(l):  
    while 1:  
        if CAS(&l->n, 0, -1)  
            return
```



RWLock: limitations

- High cache invalidation cost
 - Successful reader lock acquisition requires a cache invalidation for all other cores.
 - Does not scale well with the number of cores.
- Can we avoid this cost?
 - If the readers does not write to the shared integer, can there be speed up?

RCU: API

Readers:

- `rcu_read_lock()` // enter rcu critical section
- `rcu_read_unlock()` // leave rcu critical section

Writers:

- `synchronize_rcu()`

RCU Idea 1: Publishing Protocol

- Writer not allowed to directly modify the linked list instead, it will
 - read the current data('s next pointer),
 - copy this information to the new node,
 - and update the current data('s next pointer)
- Only works for data structures that can be atomically updated via one write
 - So doubly linked list is not a good data structure for RCU.
 - Tree is a good data structure
 - Not always a pointer: <https://www.youtube.com/watch?v=rxQ5K9lo034@14:01>
 - something if you have it gives you access to the rest of the data
 - reader gets it, gives it access to whatever data it reveals.
 - Some readers sees the old data, others see the new data; very common.

RCU Idea 2: Memory Barrier

- There is no “after this(commiting write) then”
 - Compiler & CPU will reorder instructions
 - Use memory barrier to ensure the “then” semantic
 - **Because we did not use a lock**(or other sync mechanism)
- For writer
 - the barrier goes before the “commiting write”
- For readers
 - the barrier goes after following the pointer
 - so later read from the node reads the correct(non-cached) value

RCU Idea 3: Delayed Memory Reclaim

- This is the protocol(collaboration between reader/writer) that makes RCU
- Readers
 - announce access by calling `reader_lock`
 - promise not to yield CPU in RCU critical session
 - because context switch is dual purposed to signal a reader is done accessing the shared data structure(`reader_unlock`)
 - Will always get things via this root token. If used up, traverse via the new root token up to the same location.
- Writer
 - delay free until every call has context switched at least once → by calling `rcu_synchronize` to start waiting

```
// list reader:

rcu_read_lock()
e = head
while(p){
    e = rcu_dereference(e)
    look at e->x ...
    e = e->next
}
rcu_read_unlock()

// replace the first list element:

acquire(lock)
old = head
e = alloc()
e->x = ...
e->next = head->next
rcu_assign_pointer(&head, e)
release(lock)

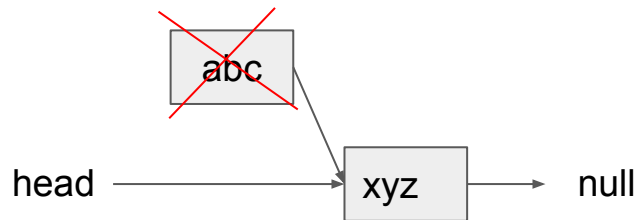
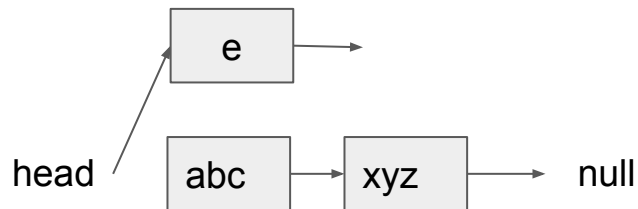
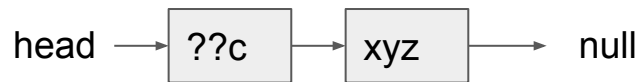
synchronize_rcu()
free(old)
```

```
void rcu_read_lock()
{
    preempt_disable[cpu_id()]++;
}

void rcu_read_unlock()
{
    preempt_disable[cpu_id()]--;
}

void synchronize_rcu(void)
{
    for_each_cpu(int cpu)
        run_on(cpu);
}
```

Figure 2: A simplified version of the Linux RCU implementation.



References

1. [CppCon 2017: Fedor Pikus “Read, Copy, Update, then what? RCU for non-kernel programmers”](#)
2. [6.S081 Fall 2020 Lecture 23: RCU - YouTube](#)